

WHAT ARE AGENT-MEMORY GAINS MADE OF? FACTORIALY DECOMPOSING REPLAY, STRUCTURAL TRANSFER, AND SURFACE TRAPS WITH EVALUATOR-ONLY PRO- GRAM ORACLES

Anonymous authors

Paper under double-blind review

ABSTRACT

Agent memory systems report consistent aggregate improvements from storing and reusing past episodes, but aggregate accuracy cannot tell *what* the gain is made of: exact replay of seen episodes, reuse of a transferable procedure under surface change, or surface-matched exploitation of an incompatible program. We decompose these channels with a benchmark of latent-program task families in which program match (P) and surface similarity (S) are independently randomized per memory–target pair while family, program, and transformation labels remain evaluator-only. Across 7,680 fixed-injection rollouts of two frozen models and a pre-registered validity audit (strong leakage probes, difficulty equivalence, continuous-similarity sensitivity, oracle re-execution), we find: (i) aggregate gains are replay-dominated—in our 7B setting 72% of the matched-content effect comes from the surface-matched leg—and the clean structural-transfer term changes sign with model scale ($\tau_{\text{struct}} = +0.092$ [Holm-significant] for 7B vs. negative for 3B with a significant $P \times S$ interaction), revising the “weaker models benefit more” conclusion drawn from system-level aggregates; (ii) near-miss memories cause paired harmful flips (3B: 18.8%; 7B: 9.2%) concentrated in states where the vis-à-vis programs semantically diverge; and (iii) in a fresh 2-by-2 factorial separating representation *form* from semantic *coverage* (9216 rollouts on 32 held-out families), replay realization is governed by coverage completeness—preserving the decision and termination steps—not by representation form: even full transcripts realize replay materially (TOST rejects a $\pm 3\text{pp}$ equivalence), necessarily revising an earlier negative result we report. Method-comparison baselines (an architecture $\times$ content transplant baseline, an embedding-based structural signal, and an LLM intent-mismatch judge) fail to recover these quantities in our harness. Deploying the coverage finding as a deterministic decision-aware compression policy beats naive top-truncation at exactly paired per-card token budgets on the pooled grid (+14.5pp replay realization at 7B, Holm-significant; at 3B there is no pooled evidence of benefit, point estimate -8.4pp , which we report with the same prominence) without degrading the primary harmful-flip guardrail at either scale (the secondary trap-shift bound is established only at 7B); a provenance-stratified sensitivity then shows the pooled 7B advantage is carried by oracle-reconstructed source episodes (direct interaction +0.38, CI excluding zero) and is not established on authentic model trajectories (+4.1pp, CI crossing zero), and we bound the deployment claim accordingly (12,800 rollouts). We release all generators, sealed-oracle evaluators, rollouts, and analysis code.

1 INTRODUCTION

Language agents increasingly improve at deployment time by reading from an external store of *memories* distilled from past episodes—verbatim trajectories, distilled reflections, procedure cards, learned workflows (Shinn et al., 2023; Zhao et al., 2024; Wang et al., 2025; Fang et al., 2025; Liang

et al., 2026; Cao et al., 2026). Reported gains are usually computed as an *aggregate* effect: task-level success with memory minus success without it. Behind a single such number, however, at least three radically different causal channels can coexist:

1. **Exact or near-exact replay.** The current task reproduces something the memory literally described, so the gain is recall of an instance, not learning.
2. **Structural (procedural) transfer.** A procedure genuinely crosses surface change—different entities, wording, domain vocabulary—because the memory encodes a reusable program.
3. **Surface exploitation and surface traps.** Tasks whose surface cues match receive (mis)guidance; when the cues are similar but the required program is different, the memory can actively hurt—a harmful flip of a task the agent would otherwise solve.

No aggregate metric distinguishes these channels, yet they carry opposite deployment meanings: replay justifies a cache, structural transfer justifies learning, traps justify admission control. Worse, the “memory effect positive / weak models benefit more” narratives (Feng et al., 2026b; Anonymous, 2025c; 2026n) are all stated at the aggregate level, and Anonymous (2026i;c) have already warned that pipeline and counterfactual confounds can fabricate or erase the conclusions.

Question. *Under identical model, budget, retriever, and prompt, what is the observed memory gain made of?* Precisely: how much of the improvement over no-memory, per memory–target pair, comes from (i) exact replay-like reuse, (ii) clean structural transfer under surface change, (iii) context/format overhead, and (iv) near-miss harm? And which of those components survives a pre-registered separation of representation *form* from semantic *coverage*?

Answer. We introduce *CausalMemBench*, a generator of latent-program task families (four confirmation archetypes rendered into four business domains), a sealed oracle that holds the program-equivalence labels evaluator-only, and a $P \times S$ factorial design (Stage A) that independently randomizes latent program match $P \in \{0, 1\}$ and surface similarity $S \in \{0, 1\}$ for each memory–target pair. The cells read: A00 (unrelated control), A01 (surface trap / near-miss), A10 (clean structural transfer), A11 (replay-like reuse), plus no-memory (N) and sham-memory (Q) references. We run a 40-family mini-pilot with frozen Qwen2.5-3B/7B, a 200-token-budget-matched injection, and family-cluster inference with pre-registered Holm correction. Then a fresh-data 2-by-2 factorial (H3) with 32 new families separates representation form from semantic coverage. All validity gates (strong leakage probes, no-memory difficulty equivalence TOST, continuous-similarity checks, oracle re-execution, blind-feature separation) pass.

Findings (each pre-registered, all reported with family-cluster CIs).

1. **Aggregate gains are replay-dominated.** In the 7B model, A11 exceeds no-memory by +18.6pp (95% CI [11.4, 26.1]), while A10—the only cell that isolates cross-surface structural reuse—contributes only a fraction of the total (72% of the pooled matched-content effect comes through the A11 leg, per analysis §5).
2. **Structural transfer flips sign with model scale.** $\tau_{\text{struct}} = A10 - A00 = +0.092$ for 7B (Holm-significant; +0.105 after difficulty adjustment), but $\tau_{\text{struct}} = -0.066$ for 3B (raw bootstrap interval excludes zero before Holm; $\tau_{P \times S} = +0.113$ 3B Holm-significant). The reading of “weak models benefit more from memory” (Feng et al., 2026b) must be revised: weak models profit from replay, not from structure.
3. **harmful flips are real and structured.** Paired harmful flips from near-miss memory reach 18.8% (3B) and 9.2% (7B), and concentrate in states where the memory’s program and the task’s program actually diverge—not uniform interference.
4. **The program-recognition assumption is weaker than believed.** A 7B judge, even when explicitly told to ignore domain nouns, cannot separate cross-domain program-equivalent cards at 0.25 agreement with generator labels (CoT), and the alarm-style intent judge saturates (AUC = 0.508).
5. **Coverage, not form, bounds replay realization.** Transcript-vs-script form contrasts are null under Holm across 8 tests; full transcripts realize replay materially ($\tau_{\text{replay}} = +0.104 / +0.190$ in 7B/3B; TOST rejects $\pm 3\text{pp}$ equivalence), while a 300-token hard truncation drops it to ≈ 0 —diagnosing our own earlier H-C anomaly as coverage, not representation.

Baselines cannot recover this. We reproduce an architecture $\times$ content transplant 2×2 (Feng et al., 2026b) inside the same harness: it measures the pooled matched effect but cannot decompose replay from structure (the 72% replay share is invisible to it; interaction -0.129 , Holm-significant). A Proc-Mem-style embedding similarity (Anonymous, 2025c) partially predicts uplift in one model (Holm-limited) but predicts neither the structural plane nor harmful flips. A STITCH-style LLM intent judge (Anonymous, 2026n) alarms on $> 90\%$ of cells of all kinds (AUC=0.508).

Honest scope. We report two formal negative results under pre-registered gates: a write-path form comparison (raw/summary/procedural) that did not pass its minimal gate, and the H3 factorial’s frozen criterion (form effect), reported verbatim with mechanism-level clarification. TRU-Mem, a transportable-utility admission method, was deliberately *not* submitted for review: its advertised differentiation is not up to its own pre-registered standard (see §9).

Contributions. (1) The first pair-level $P \times S$ orthogonal identification with evaluator-only latent-program oracle for agent memory, with a complete audit protocol. (2) Empirical corrections to aggregate memory narratives: replay dominance, scale reversal of structural transfer, structured harmful flips, weak program equivalence of LLM self-judgment, and the coverage-not-form law of replay realization. (3) Baseline-in-harness evidence that transplant-factorial, embedding-based, and LLM-judge signals do not predict these quantities. All task generators, sealed evaluators, 28,416 rollouts, probes, blind annotations, and analysis scripts are public.

2 RELATED WORK

Procedural and episodic agent memory. Reflexion stores verbal reflections for iterative self-correction (Shinn et al., 2023); ExpeL learns natural-language insights across episodes (Zhao et al., 2024); Agent Workflow Memory induces reusable workflows from web trajectories (Wang et al., 2025); Memp systematizes procedural-memory build/retrieve/update (Fang et al., 2025); MCMA adds multi-level abstraction meta-learning (Liang et al., 2026); ReMe couples distillation with utility-based refinement (Cao et al., 2026). Benchmarks characterize incremental memory use (Anonymous, 2025b;a; 2026a;g). These measured *whether* memory helps, in aggregate, and largely leave its *what* unresolved. Audit-style works point at confounds: Anonymous (2026i) on pipeline factorization, Anonymous (2026c) on static-vs-causal utility, Anonymous (2026h) on poisoned-memory attribution, and Anonymous (2026e) on consolidation decay.

Factorial and transfer audits. Memory Transplants (Feng et al., 2026b) is the closest precedent: a 2-by-2 factorial of memory architecture and stored content over a code-to-math shift, with prompt freeze and token controls, at a *system/domain* level. Operation-level follow-ups vary retrieve/write/manage (Feng et al., 2026c); AIM models interference dynamics at retrieval boundaries (Feng et al., 2026a). We run the same kind of transplant-in-harness comparison (Section 8): the design cannot, by itself, separate the surface-matched leg from the cross-surface leg of a matched-content gain. Rocchi’s SSRN preprint audits single memories at bank $\times$ decoder granularity with outcome oracles (Rocchi, 2026); our axes and labels are complementarily defined (program-match, surface similarity, latent-equivalence oracle). Coding-domain cross-domain analyses show abstract memories travel better than low-level traces but can negatively transfer (Anonymous, 2026j); skill libraries suffer shadowing (Anonymous, 2026l).

Similarity, retrieval, intent. Proc-Mem reports an embedding “generalization cliff” on novel procedure vocabulary (Anonymous, 2025c); STITCH/CAME-Bench studies semantically-similar-but-intent-mismatched memories and gates them by latent-goal/action/entity types (Anonymous, 2026n). We test both as predictors (§8): neither embedding similarity nor LLM intent judgment recovers our causal quantities.

Admission, utility, risk control. Utility-factor admission (Zhang et al., 2026), risk-sensitive abstention (Anonymous, 2026m), decision-aware utility scoring (Anonymous, 2026d), write-time consistency (Zhang & Li, 2026), transactional commit (Anonymous, 2026k), hindsight or attribution-based credit (Anonymous, 2026f;b), and conformance certificates for memory-grounded actions (Akewar & Ranjan, 2026) are method-side designs that consume utility estimands. Our framework is the measurement layer that validates those estimands; we do not compete with their admission rule. Risk-guarantee machinery (Angelopoulos et al., 2024; Wu et al., 2016) and transportability theory

(Jalaldoust & Bareinboim, 2024) inform our framing, but we make no risk-control theorem claim in this paper.

Position. To our knowledge no prior work independently randomizes latent program match and surface similarity per memory–target pair, holds program equivalence evaluator-only, and estimates the decomposition with family-cluster inference plus a full audit chain (leakage, difficulty, similarity calibration, oracle re-execution, blind features).

3 IDENTIFICATION: WHAT THE DECOMPOSITION NEEDS

3.1 WHY AGGREGATE EFFECTS ARE NOT IDENTIFIED

Let $Y(x, W)$ be the success indicator of an agent at task x with memory bank W , and let m be one stored episode. An aggregate contrast $\Delta_{\text{agg}} = \mathbb{E}[Y(x, \{m\}) - Y(x, \emptyset)]$ is compatible with at least two distinct data-generating processes: (i) an *exact-replay* SCM where Y improves only when x literally reproduces m ’s instance; (ii) a *structural-transfer* SCM where Y improves whenever x shares m ’s program, however re-surfaced. Any estimator based on aggregate differences alone cannot distinguish them; we need interventions on the relationship between m and x itself (§3.2).

3.2 LATENT PROGRAM EQUIVALENCE

A latent program $z = (G_{\text{prec}}, C_{\text{safety}}, C_{\text{terminal}}, B_{\text{recovery}})$ specifies partial-order dependencies G over abstract steps (READ, CHECK, WRITE), safety rules (writes gated on conditions), legal terminal predicates over the environment state, and minimal recovery requirements. The equivalence class $\Pi(z) = \{\pi : \pi \models z\}$ collects all action sequences satisfying these constraints *irrespective of* which of several valid linear orders is taken, which entities appear, or which business domain renders the task. Two tasks are program-matched ($P=1$) iff they belong to the same class—independent of vocabulary, entity names, narrative style, and surface tool names. Near-miss z' is a deliberate single mutation: flipped comparison polarity, reversed transfer direction, wrong child-set aggregation, or a missing archival step; it is executable and reaches its own legal terminal state.

3.3 THE $P \times S$ STAGE-A FACTORIAL

For every target sibling x and its source episode, we generate five memory conditions (length-matched cards):

Cell	(P, S)	Content
A00	(0, 0)	correct procedure for an unrelated program/domain (control)
A01	(0, 1)	near-miss: surface-similar, program-different (trap)
A10	(1, 0)	same program, different entities/vocabulary (clean structural)
A11	(1, 1)	same program, similar surface (replay-like)
Q	—	sham: format/length/success-tag matched, task-irrelevant content

plus N (no memory). P is randomized by latent family; S buckets are calibrated by two frozen similarity views (token-overlap TF-cosine on content tokens; bge-small embedding cosine), and P, S are statistically independent within family so their marginal and interaction effects are factorially identified under consistency, positivity (all cells populated per family), and no cross-family interference (fresh deep-copied environment state per rollout).

Denote by $Y(p, s)$ the success rate with memory of type (p, s) , by Y_N, Y_Q the no-memory and sham rates. Pre-registered estimands, all in risk-difference scale: context/format effect $\tau_{\text{context}} = \mathbb{E}[Y_Q - Y_N]$; program main effect $\tau_P = \frac{1}{2} \sum_s \mathbb{E}[Y(1, s) - Y(0, s)]$; **clean structural transfer** $\tau_{\text{struct}} = \mathbb{E}[Y(1, 0) - Y(0, 0)]$; **surface-trap** $\tau_{\text{trap}} = \mathbb{E}[Y(0, 1) - Y(0, 0)]$; replay-like premium $\tau_{\text{replay}} = \mathbb{E}[Y(1, 1) - Y(1, 0)]$; interaction $\tau_{P \times S} = (\mathbb{E}[Y(1, 1) - Y(0, 1)]) - (\mathbb{E}[Y(1, 0) - Y(0, 0)])$; paired harmful-flip rate $\text{HFR} = \Pr[Y_N=1 \wedge Y_{A01}=0]$ on matched seeds.

Estimation uses family-cluster bootstrap (10,000 / 20,000 resamples; percentile 95% CIs) with Holm multiplicity control across the four primary estimands per model (Section 4).

3.4 EVALUATOR-ONLY ORACLE

The oracle holds: family index, abstract signature (the canonical equivalence key), program parameters, near-miss label, cell assignments, generator seed. Only the evaluator reads it; the agent sees task text, tool schema, and current state. This blocks four leakage paths: (i) cards contain no family/cell IDs (grep-enforced); (ii) generators and rollouts are seed-logged; (iii) style templates are Latin-square balanced across cells; (iv) evaluation utilities never read sealed fields—leading to a sealed-evaluator architecture of benchmark release (Section 4, audit).

3.5 FROM PILOT TO FORM-COVERAGE FACTORIAL (H3)

A minimal-gate experiment on three write-path representation forms (verbatim transcript, summary, procedural card, §6) produced a pre-registered outcome we must account for: a large replay-premium difference ($\Delta_{\tau_{\text{replay}}} = -0.158$), but the raw cards were systematically truncated at the 300-token budget, confounding *form* with *coverage*. We therefore pre-registered H3 on 32 *fresh, disjoint* families: a 2×2 factorial of **form** (dialogue-transcript vs. imperative script) $\times$ **coverage** (complete incl. decision+finish vs. matched prefix without them), constructed from one canonical proposition set per episode (100% programmatic content alignment across forms), plus one ecological arm—the verbatim H-C 300-token hard cut. H3’s frozen criteria are the form contrast $\varepsilon_{\text{form}}$, per-form coverage contrasts ε_{cov} , and their interaction, with TOST on transcript-complete τ_{replay} at margin $\pm 3\text{pp}$.

4 BENCHMARK, HARNESS, AND VALIDITY AUDIT

4.1 ENVIRONMENT AND TASK FAMILIES

RelationalOps is a SQLite-backed tool environment with four business templates (CRM, inventory, ticket, calendar) and a six-tool protocol (`list_tables`, `read`, `aggregate`, `insert`, `update`, `delete`, `finish`). Success is judged by *programmatic terminal predicates* over the resulting database state—no LLM judge anywhere. We instantiate four partial-order archetypes: (P1) conditional write (`READ`; `CHECK`; `BRANCHWRITE`; `VERIFY`), (P2) two-row transfer with conservation guard, (P3) aggregate gate with audit log, (P4) delete-after-capture (archive then delete), each rendered in two domains, producing 40 pilot latent-program families and 32 fresh H3 families (disjoint generator seeds; 800 and 640 tasks respectively).

4.2 HARNESS

Agent episodes use frozen backbones (Qwen2.5-3B, Qwen2.5-7B; Team, 2024) under vLLM temperature 0.7, top- p 0.9, up to 12 steps of 512 tokens each. For every (sibling, seed) pair, all six cells reuse a byte-identical initial DB and the same decoding seed; only the injected memory differs. Cards are 200–300 Qwen2.5-1.5B tokens, n-style-templated, Latin-square balanced. Parseable-action rates: 98.9% (3B), 99.6% (7B); compliance (memory actually consulted) verified on three layers (step-level coverage, action-sequence deltas vs. N, and trace-level echo).

4.3 THE VALIDITY AUDIT THAT ACTUALLY GATES THE CLAIMS

All claims below passed, in advance of analysis:

1. **Oracle executability.** 800/800 siblings walk to legal terminal states; 160/160 re-executed in audit.
2. **Structural alignment.** The two-domain renderings of each family share byte-identical abstract signatures (8/8 exact), and every near-miss differs by exactly one designed mutation.
3. **Leakage probes at high strength.** In the structural plane (A00 vs. A10), char n-gram TF-IDF AUC = 0.378 [0.150, 0.624], length/style AUC = 0.346, embedding AUC = 0.380, LLM probe AUC=0.724—none able to recover treatment from public text at held-out families; family-index 40-way classification accuracy = 0.087 (chance 0.025).

Table 1: Stage-A success rates with 95% family-cluster CIs (n=640/cell/model). Bold marks the replay-dominated leg.

Model	N	Q	A00 (0, 0)	A01 (0, 1)	A10 (1, 0)	A11 (1, 1)
Qwen2.5-3B	.420	.402	.373	.362	.308	.409
Qwen2.5-7B	.547	.573	.497	.578	.589	.733

Table 2: Pre-registered Stage-A estimands (risk differences; Holm over 4 primary estimands per model). Starred = Holm-significant.

Estimand	Qwen2.5-3B	Qwen2.5-7B
$\tau_{\text{context}} = Q - N$	-0.019 [-0.067, 0.031]	+0.027 [-0.023, 0.080]
$\tau_{\text{struct}} = A10 - A00$	-0.066 [-0.131, -0.002]	+0.092 [+0.036, +0.152]*
$\tau_{\text{trap}} = A01 - A00$	-0.011 [-0.072, +0.047]	+0.081 [-0.006, +0.173]
$\tau_{\text{replay}} = A11 - A10$	+0.102 [+0.039, +0.161]*	+0.144 [+0.087, +0.200]*
$\tau_{P \times S}$	+0.113 [+0.041, +0.189]*	+0.063 [-0.031, +0.156]
HFR _{A01}	18.8% [15.1, 22.5]	9.2% [6.3, 12.0]

4. **Difficulty equivalence.** No-memory sibling-difficulty equivalence by TOST at $\pm 7\text{pp}$: 3B passes directly; 7B passes after trimming four worst-matched pairs (estimands recomputed post-trim get *stronger*: $\tau_{\text{struct}} = +0.105$).
5. **Continuous-similarity sensitivity.** Binary- S conclusions reproduce as bucket-boundary shifts on continuous sim (tf/embedding); slopes within matched strata are flat where bins were defined.
6. **Identification advantage.** Saturating on continuous sim measures, the structural coefficient still holds at +0.165 (7B)—architecture/content dummies of the transplant 2×2 cannot absorb it.

4.4 DATA RECORDS

28,416 rollouts (7,680 pilot + 7,680 H-C + 9,216 H3 + 3,840 Llama-8B replication) with per-row git commit, config hash, seeds, GPU model, and vLLM/torch versions; all failures logged as-is; no rollout fabricated or fabricable (Section B).

5 FINDINGS

5.1 AGGREGATE GAINS ARE REPLAY-DOMINATED

Table 1 lists Stage-A cell success rates (n=640 per cell per model). For Qwen2.5-7B, no-memory solves 54.7%; replay-like A11 reaches 73.3% (+18.6pp over N, CI [11.4, 26.1]), while structural A10 reaches 58.9% and context Q is statistically indistinguishable from N ($\tau_{\text{context}} \approx 0$, CI spans 0). For 3B, A11 is the only memory cell that stays at parity with N ($\tau_{\text{replay}} = +0.102$ significant over A10), while every other injected memory degrades performance below N.

Decomposing the pooled *matched-content* effect ($A11 \cup A10$ minus A00) yields 72% attributable to the A11 (replay-like) leg for 7B; the pooled τ_P is therefore a poor summary—the decomposition is the object of interest.

5.2 STRUCTURAL TRANSFER FLIPS SIGN WITH MODEL SCALE

The structural channel is real for 7B (τ_{struct} Holm-significant; +0.105 after difficulty-adjusted trimming, still significant) and inverts at 3B: the interaction $\tau_{P \times S}$ tells us the 3B’s aggregate advantage is driven by surface match alone ($S=1$, P only helps then). A transplant-style aggregate read of the same data (Feng et al., 2026b) reports the mixture “memory helps less at larger scale”—but the decomposition says: *it is structural transfer that scales, not helpfulness*. Weak models profit from *replay* (3B τ_{replay} significant) and are unprofitably steered by structurally-correct-but-foreign-surface

Table 3: Replay-related cell rates and contrasts on fresh families (n=384/arm-cell/model). Arms: S/T=script/transcript; C/P=complete/prefix; eco = verbatim 300-token hard cut.

arm	Qwen2.5-7B		Qwen2.5-3B	
	A10	A11	A10	A11
T-complete	.729	.833	.312	.503
S-complete	.661	.732	.299	.424
T-prefix	.568	.591	.224	.312
S-prefix	.526	.544	.242	.273
eco	.641	.638	.245	.297

guidance. This is our first revision to a published aggregate conclusion (§8 checks baselines cannot substitute this inference). A second-family pilot on Llama-3.1-8B (Appendix D) reproduces the same ordering—replay dominating (79% of the matched effect from A11) with a 7B-like positive structural term—arguing the sign pattern tracks capability rather than a Qwen idiosyncrasy.

5.3 NEAR-MISS HARM IS STRUCTURED, NOXIOUS—AND MODEL-SPECIFIC

Paired harmful flips (so that $Y_N=1$ while $Y_{A01}=0$ on matched seeds) are large: 18.8% for 3B [15.1, 22.5] and 9.2% for 7B [6.3, 12.0]. The marginal τ_{trap} is deceptively small because helpful and harmful flips partially cancel; the flip metric exposes what the mean hides. Crucially, the harm is not a mist: on a per-archetype-by-branch decomposition it concentrates exactly where the near-miss and correct programs *semantically diverge*. For example on our P3 archetype (gate-checking over the wrong child set) harmful flips are 17.5–32.5% on the branch where the two programs disagree, vs [0, 12.5]% where they coincide. The trap is carrying information, and the model is failing to discount it where the implicit program is wrong (Section 5.4 for why models cannot self-gate this harm).

5.4 MODELS CANNOT EVEN RECOGNIZE PROGRAM EQUIVALENCE

An independent 7B judge asked to verify, for the same (task, memory) pair, whether the procedures are program-equivalent—with explicit mention to ignore domain nouns—agrees with generator labels at only 0.250 (CoT) over 32 blind judgments. It calls every cross-domain program-equivalent pair (A10) “different” (100% error on A10) and misses all direction-reversal and child-set near-misses (0/4 each). An intent-mismatch judge in the style of Anonymous (2026n) fires an alarm on > 90% of cells of all kinds (detector AUC=0.508 vs. chance). The practical implication: the system that consumes a memory cannot verify the memory by text-at-scale, which is exactly why evaluator-side randomization is needed to diagnose it in the first place (§3).

6 REPLAY REALIZATION IS BOUNDED BY COVERAGE, NOT FORM

6.1 AN ANOMALY WE HAD TO EXPLAIN

A pre-registered H-C minimal gate compared three write-path forms of the same source episodes under token-parity (raw transcript truncated to ≤ 300 tokens; LLM summary; procedural card) §3.5. The celebrated “script-over-raw” gap that at first read suggested form effects was large ($\Delta\tau_{\text{replay}} = -0.158 [-0.239, -0.080]$)—but every raw card had been hard-truncated and systematically omitted the decision step (only 48–59% contained a write step; 8 of 640 contained the finish step). The “form” claim was unidentified.

6.2 SEPARATING FORM FROM COVERAGE

With one canonical set of propositions per episode (100% programmatically content-aligned across forms), the fresh 32-family slice renders clean cells: transcript-vs-script complete/prefix, plus the verbatim H-C-style 300-token hard cut as an ecological arm. Success rates (4320 arm roll-outs/model):

Cell-level replay terms $\tau_{\text{replay}} = A11 - A10$: transcript-complete +0.104 ([+0.042, +0.171], 7B) and +0.190 ([+0.123, +0.259], 3B); script-complete +0.070 and +0.125; the form contrast $\varepsilon_{\text{form}} = \tau_{\text{replay}}^{(\text{script})} - \tau_{\text{replay}}^{(\text{transcript})}$ is negative in both models (Holm over 8 contrast tests: all null). Coverage contrasts are directionally positive in both forms and models (raw 3B transcript $p = .018$) but do not survive the frozen family-wise correction. TOST on transcript-complete τ_{replay} at margin $\pm 3\text{pp}$ is rejected in both models: *full transcripts realize replay materially*. The eco arm, forced to the same 300-token brutality as H-C, collapses to ≈ 0 in 7B (-0.003 [$-0.083, +0.083$]), mirroring H-C's earlier -0.014 .

6.3 RESOLUTION

Our earlier H-C anomaly is thus identified as *sensitive to coverage/truncation*, not to representation. We deliberately state its strength precisely: the result shows transcripts *can* realize replay when their decision content is preserved; it does not isolate a clean mechanism status for coverage (multiplicity-corrected coverage contrasts are not significant). See Section 9 for the mapping to the overall strength of evidence.

7 DEPLOYING THE FINDING: DECISION-AWARE COMPRESSION

Section 6 ended on a constructive statement: transcripts realize replay *if* their decision content survives. We now test whether that statement can be turned into a compression *policy* that beats the naive baseline an engineer would actually ship: top-truncation to the token cap.

7.1 A PAIRED-BUDGET COMPRESSION TEST

The policy (DSLICE) is deterministic—no LLM, no paraphrase, no generated natural language and no new semantic values. The *base card* (rules D1–D5) retains one canonical copy of each distinct action (exact-duplicate non-mutating actions are collapsed); compresses each read/aggregate tool result to a one-line summary (table, filter, matched-row count, retrieval-key columns of the first matched row); drops list-table dumps and write-operation results; and keeps the finish step and its result. Cards still exceeding 300 tokens escalate deletion in a fixed order: of 640 cards, 14 fit at the base stage and the rest delete successively read-summary key columns (76 cards), read result lines (281), the finish *result* (8), non-decision action lines (252), and aggregate action lines (9). What is retained at 100% by assertion is therefore narrower than the base card: parsed write (insert/update/delete) action arguments, every aggregate value, and the finish *action* (0 QA failures over 640 cards; construction audit in Figure 5). Cards average 237.5 tokens (range 102–299).

The control is the engineer's baseline done fairly: the *same* full transcript, top-truncated not to 300 tokens but to *exactly the token count of the paired* DSLICE card (RAW_MATCHED; per-card $|\Delta\text{tokens}| = 0$ for 640/640 pairs). Length can therefore not confound the comparison in either direction. One property of the source deserves emphasis: both card types inherit the provenance of the H-C raw harvest, and for 188 of the 640 memories (29.4%; 52/49/43/44 in A00/A01/A10/A11) no harvest attempt succeeded, so the stored trajectory is a deterministic reconstruction of the evaluator's oracle plan rather than a model-generated episode. All arms consume the identical 640 sources, so pairing and budget exactness are unaffected; but the pooled contrast mixes two provenance strata, which we stratify post-hoc in §7.3. We also report the ecological 300-token cut (RAW(300)) as the deployment senior baseline. A timestamped pre-analysis protocol and its amendments (including the addition of the 3B arm and the paired-budget control; registration provenance disclosed in Appendix A) govern the design; 12,800 new rollouts cover the full grid $\{40 \text{ families} \times 4 \text{ siblings} \times 4 \text{ cells} \times 4 \text{ seeds} \times 2 \text{ models}\}$ for DSLICE and RAW_MATCHED, plus the missing RAW(300) 3B arm.

7.2 RESULT: A LARGE HOLM-SIGNIFICANT REALIZATION GAP AT 7B, A POST-HOC POOLED SCALE DIFFERENCE

Table 4 (rates and CIs visualized in Figure 6) is the pooled evidence. On 7B, naive truncation at *exactly* the same per-card budget realizes no replay at all ($\tau_{\text{replay}} \approx -0.03$), while DSLICE realizes +0.112; the paired gap $E1 = +14.5\text{pp}$ is Holm-significant. The guardrails behave as required: HFR is, if anything, *lower* under DSLICE on both models (noninferiority established at the +5pp

Table 4: Decision-aware compression vs. paired naive truncation (family-cluster bootstrap, 2000 reps). $\tau_{\text{replay}} = A11 - A10$; E1 is the paired difference of replay terms (DSLICE minus RAW_MATCHED); E2 is the DSLICE replay term itself. All p -values in this table are Holm-adjusted within the primary $m = 6$ family. Guardrails are one-sided noninferiority checks: ΔHFR (upper-95 $< +5\text{pp}$) is primary, a member of that family; $\Delta\tau_{\text{trap}}$ (lower-95 $> -5\text{pp}$) is secondary, outside the Holm family (“not est.” = bound crosses the margin). DSLICE cards: mean 237.5 tokens.

arm	Qwen2.5-7B		Qwen2.5-3B	
	A10	A11	A10	A11
DSLICE	.558	.670	.397	.438
RAW_MATCHED	.559	.527	.242	.367
RAW(300)	.598	.584	.225	.389
procedural (ref.)	.589	.733	.308	.409

contrast	7B	3B
E1 (paired $\Delta\tau_{\text{replay}}$)	+0.145 [+0.048, +0.233], $p = .012$	−0.084 [−0.177, +0.002], $p = .97$
E2 ($\tau_{\text{replay}}^{\text{DSLICE}}$)	+0.112, $p = .051$	+0.041, $p = .66$
ΔHFR (primary)	pass (u95 −0.016)	pass (u95 +0.011)
$\Delta\tau_{\text{trap}}$ (secondary)	pass (+0.139; lower 95 +0.075)	not est. (+0.000; lower 95 −0.072)

margin), and at 7B the pair shows a secondary τ_{trap} shift of +0.139 in the protective direction (lower 95 +0.075). At 3B τ_{trap} is unchanged on average (+0.000) and the −5pp noninferiority bound is not established (lower 95 −0.072)—we report this rather than pass it. At 3B we otherwise find no pooled evidence of benefit: the E1 point estimate is −8.4pp against RAW_MATCHED (the CI [−0.177, +0.002] excludes neither a moderate harm nor zero), and the RAW(300) comparison is negative (secondary, unadjusted CI −0.123 [−0.216, −0.039]). A post-hoc cross-model difference of E1 (+0.230 [+0.131, +0.338], outside the Holm family) supports a pooled effect that differs by scale; it speaks to the size of the pooled effect, not to a capability mechanism—and it is subject to the provenance caveat of §7.3, which applies to both scales.

7.3 STRENGTH OF EVIDENCE

What the pooled result licenses. A *packaged* deterministic policy derived from the coverage finding beats naive truncation at an exactly paired budget on the larger model, with the primary harmful-flip guardrail intact at both scales and the secondary trap-shift guardrail established at 7B. We do not claim that any single rule inside the package (action retention, result summarization, escalation order) is responsible, and DSLICE does not reach procedural-card realization (A11 .670 vs. .733): decision-aware compression of a transcript closes much of the realization gap but does not close it.

What it does not license. A post-hoc provenance stratification (the 188/640 oracle-reconstructed sources of §7.1; family-cluster bootstrap over the same resample semantics as the primary analysis; outside the Holm family; numbers in Appendix A) materially qualifies the interpretation. On authentic model-generated trajectories the 7B advantage is +0.041 [−0.078, +0.149], while on oracle-reconstructed trajectories it is +0.424 [+0.255, +0.615]; at 3B the strata give −0.226 [−0.347, −0.104] and +0.297 [+0.167, +0.473]. A direct provenance interaction (one aligned family-cluster bootstrap over the union family universe) converts this from two separate stratum impressions into one heterogeneity estimate: oracle-minus-authentic E1 is +0.382 [+0.166, +0.606] at 7B and +0.523 [+0.306, +0.783] at 3B. We take no position on the mechanism: the DSLICE replay term is nearly identical across strata (+0.117 authentic vs. +0.104 oracle at 7B) while the RAW_MATCHED control swings (+0.076 vs. −0.319)—the strata differ mostly in how much the naive baseline loses, for reasons this design does not identify. We therefore state the deployment claim precisely: DSLICE-style compression beats exact-budget truncation as a *benchmark* contrast on this mixed-provenance grid, with no pooled evidence of benefit at 3B (point estimate −8.4pp); whether the policy helps when the stored episodes are genuinely model-written remains open, and answering it requires a fallback-free re-harvest (new rollouts) that we flag as the blocking follow-up experiment rather than assume away. The pre-analysis amendments that make the pooled compar-

ison fair (paired-budget control, 3B completion, tokenization no-op verification) are disclosed in Appendix A.

8 BASELINES CANNOT RECOVER THE DECOMPOSITION

We ask whether existing evaluation signals substitute for URN oracle-randomized decomposition, by embedding three strongest neighbors *inside our own harness* with identical models and budgets.

8.1 THE TRANSPLANT 2×2 (FENG ET AL., 2026B) CANNOT TELL REPLAY FROM STRUCTURE

We reproduce their architecture $\times$ content factorial structure: two write-path architectures (procedural card; raw transcript) $\times$ three content conditions (matched correct (A11 $\cup$ A10); near-miss; unrelated) on identical families (N=3840+3840 rollouts, model = Qwen2.5-7B). It reports matched-content effects (+0.164, +0.035 n.s.) but *by construction* cannot decompose the A11-vs-A10 legs of what counts as “matched”. Our $P \times S$ split is Holm-significant in both procedural models (+0.144, +0.102), is zero on raw (−0.014), and the arch $\times$ content interaction $I_{\text{matched}} = -0.129$ is also Holm-significant—which shows: (i) representation changes the replay share, not only the total; (ii) nothing in the arch $\times$ content frame tells one that in this experiment 72% of the pooled matched gain is replay. Pair-level $P \times S$ is not a higher-resolution version of transplant factorial; it asks a question the 2 by 2 cannot define.

8.2 EMBEDDING-BASED STRUCTURAL SIGNAL (ANONYMOUS, 2025C): ONLY PARTIALLY PREDICTIVE

Memory–task embedding similarity (frozen bge-small, our calibration view) reaches AUC 0.60 at predicting P within the trap plane and 0.35 – −0.38 in the structural plane (chance), plus a positive sim $\rightarrow$ uplift slope significant only for 7B after Holm correction ($p_{\text{Holm}} = 0.032$). Critically it does not predict harmful flips ($p > 0.1$, odds ratios). Distilling F-MED to an embedding index thus erases both the trap direction and the clean structural estimate itself.

8.3 LLM INTENT-MISMATCH JUDGES (ANONYMOUS, 2026N): MISCALIBRATED AT SCALE

The 512-judgment blind alarm on (task, memory) pairs fires on A00/A01/A10/A11 at 100%/93.8%/100%/90.6% (AUC=0.508, i.e., no usable discrimination), consistent with the program-equivalence failure in Section 5.4. Family-level alarm does not predict per-family harmful-flip either (n.s.).

9 WHAT WE DELIBERATELY DID NOT CLAIM

1. **H-C was a formal NO-GO.** The three-form write-path minimal gate failed its frozen criterion. We did not reframe it as a positive script-form effect; we diagnosed its truncation artefact with H3 and report both.
2. **H3 was a formal NO-GO on its primary estimand.** The frozen form criterion did not reach significance; coverage is supported (directionally consistent, two models, robust simple effects) but *not confirmed* under the frozen multiplicity family. We write it as “supported-but-unresolved”.
3. **TRU-Mem removed.** A transportable-utility admission learned against this oracle at minimal risk is the obvious applied sequel, but the method-side delta over Zhang et al. (2026); Anonymous (2026m;d); Cao et al. (2026) did not meet our own pre-registered differentiation standards. It is not submitted for review.
4. **No universal invisible-leakage assertion.** We report calibrated probe AUCs and observed strength limits, not “zero leakage proofs”.
5. **Published-system profile reversal not pursued.** Two pre-registered attempts (H-C, H3) did not find significant divergence between representation forms; we do not mount a post-hoc third hunt with increased power (which would be power-chasing).




Figure 1: Stage-A cell success rates with 95% family-cluster CIs. A11 (replay-like) dominates the 7B pooled gain; the 3B pattern differs—only A11 gets near-parity with N while A10 (clean structural) falls below the unrelated control, foreshadowing the scale reversal of Section 5.2.




Figure 2: Risk differences relative to N and Q. The format/context effect ($Q - N$) hugs zero; the decomposition separates a small structural term from a large replay term; 3B only benefits through A11.

10 DISCUSSION

10.1 WHAT CHANGES IN THE AGGREGATE NARRATIVE

Our three main numbers are uncomfortable reading for several widely-quoted results. The pooled matched-content effect that aggregate benchmarks report is replay-dominated here (72% from the A11 leg), but the same pooled statistic would have survived two very different decompositions—a trap (A01) moving 18.8% of solvable tasks to failure for 3B, or a clean 9.2pp structural transfer for 7B. “Memory helps” is an aggregate claim that the community has been shipping where it wasn’t measurable. System-level transplant experiments could not have caught this difference (§8): their factors change systems, not the semantic relation between a candidate memory and a target task.

10.2 DEPLOYMENT READING

Concrete, directly-usable implications:

1. **Stop consulting aggregate accuracy** when comparing memory systems. Report the four numbers: context, clean structural, replay, and flip rate.
2. **Preserve the decision+finish of an episode** in any stored form. Our truncation results show that a 300-token cap that silently eats the decision step turns useful experience into effectively zero-signal text ($\tau_{\text{replay}} \approx 0$), even with perfect surface match—and a deterministic decision-aware




Figure 3: Paired harmful flips, A01 cell. Concentration on branches where the near-miss program semantically differs from the correct program, not a uniform effect; magnitude falls with scale.




Figure 4: Form/coverage arms on fresh families (H3). Top pad: A10/A11 rates per arm; bottom pad: replay term per arm with 95% CIs. Complete transcripts realize replay materially (TOST rejects $\pm 3\text{pp}$ equivalence), hard-truncated (eco) does not; form contrast is Holm-null in 8 tests.

compression at *exactly* the same per-card budget realizes $\tau_{\text{replay}} = +0.112$ on a 7B worker (+14.5pp over paired truncation on the pooled grid, Holm-significant; §7). Both halves of that sentence are provisional: at 3B there is no pooled evidence of benefit (point estimate -8.4pp), and a provenance sensitivity shows the pooled 7B advantage carried by oracle-reconstructed episodes rather than authentic model trajectories (direct interaction CI excluding zero), so the prescription awaits a fallback-free re-harvest.

3. **Do not trust an LLM’s own equivalence check** for memory admission. We tested retrieval-grade judgment plus intent-mismatch alarms; neither recovers the label at usable strength in our environment. Program-equivalence estimation is an inference problem, not a lookup.
4. **Treat near-miss content as a fault channel**, not as style noise: paired harmful flips are model-scale-dependent (18.8% at 3B, 9.2% at 7B) and branch-conditional on program divergence.
5. **Expect structural transfer to be capability-gated**, not universal: the structural term changed sign between 3B and 7B under identical content and budget. If your fleet includes small models, measure whether structural memory *hurts* it before shipping.




Figure 5: DSLICE card construction audit (Section 7). Left: realized card lengths (mean 237.5 tokens vs. the 300 cap); the paired RAW_MATCHED control is cut to each card’s exact count, so token budgets match per card by construction. 188/640 sources are oracle-plan reconstructions of failed harvests (provenance disclosure in §7.1); all arms share identical sources. Right: highest deletion stage reached per card (1: read-summary key columns; 2: read result lines; 3: finish *result*; 4: non-decision action lines; 5: aggregate action lines). Asserted invariants—parsed write-action arguments, aggregate values, finish action—hold in 640/640 cards (0 QA failures).

10.3 LIMITATIONS

- 1. Synthetic environment.** All identification is by construction in one SQLite world with programmatic predicates. Our ALFWorld check (Appendix E) shows the replay/structural ordering repeats at coarse level in a non-synthetic environment but the near-miss harm was contingent on trap construction—directions replicate at aggregate level, magnitudes don’t. We still do not claim AppWorld-scale external validity.
- 2. Narrow model coverage.** Two frozen Qwen2.5 scales drive the main inference; a Llama-3.1-8B pilot replication (same 40 families, 3840 rollouts) confirms replay-dominance (A11 leg = 79% of matched effect), the positive structural signal (7B-like, $\tau_{\text{struct}} = +0.047$ [$+0.020, +0.073$]), and similar flip magnitude (HFR=0.100). Directionally consistent, but the family count remains small, and Llama shows a small positive $\tau_{\text{context}} = +0.034$ absent in Qwen: family idiosyncrasies exist.
- 3. Equivalence-class ontologies.** Our partial-order classes are stricter than unconstrained “task similarity”; alternative operationalizations (e.g., embedding-based clustering of hidden work-flows) are not evaluated and could reshuffle A10 vs A01 at the boundary.
- 4. Trap-plane surface readability.** Chip-level near-miss detection from text was possible at family hold-out (AUC 0.996); we register it as a semantics of benchmark traps rather than leakage, but it limits what “unreadable label” claims can be made on that axis.
- 5. Analysis breadth.** worry: power is limited by 40/32 family clusters; decompositions at archetype level, though reported, have wide intervals and are exploratory only.

10.4 WHAT SHOULD COME NEXT

The applied sequel is a transportable-memory admission rule validated the same way—we did not run it because it would be governed against standards we cannot yet meet (Section 9). The theory sequel is a distributional-robustness treatment of transformation-shifted memory utility (Angelopoulos et al., 2024; Wu et al., 2016; Jalaldoust & Bareinboim, 2024). And the measurement sequel asks how far these estimates agree with the observational estimators people can only compute without our oracle.

REFERENCES

Mayur Akewar and Ravi Ranjan. SafeCommit: Certifying when memory-grounded agents may safely act. arXiv:2608.04289, 2026.




Figure 6: Deployment result (Section 7). Top: A10/A11 rates per arm per model, family-cluster 95% CIs. Bottom: replay term per arm. On the pooled 7B grid, DSLICE realizes replay while exactly-paired truncation realizes none (E1 = +0.145, Holm $p = .012$); at 3B there is no pooled evidence of benefit (E1 = -0.084 [-0.177, +0.002], Holm $p = .97$). A post-hoc provenance stratification qualifies both: on authentic model-generated sources E1 = +0.041 [-0.078, +0.149] (7B) and -0.226 [-0.347, -0.104] (3B), and a direct interaction estimates oracle-minus-authentic E1 at +0.382 [+0.166, +0.606] (7B) and +0.523 [+0.306, +0.783] (3B); the pooled 7B win is carried by oracle-reconstructed sources (Appendix A).

Anastasios Angelopoulos, Stephen Bates, et al. Conformal risk control. In *ICLR*, 2024.

Anonymous. LifelongAgentBench: Evaluating LLM agents as lifelong learners. arXiv:2505.11942, 2025a.

Anonymous. Evaluating memory in LLM agents via incremental multi-turn interactions. arXiv:2507.05257, 2025b.

Anonymous. A benchmark for procedural memory retrieval in language agents. arXiv:2511.21730, 2025c.

Anonymous. Managing procedural memory in LLM agents. arXiv:2606.23127, 2026a.

Anonymous. AttriMem: Attribution-guided process feedback for agent memory learning. arXiv:2607.21106, 2026b.

Anonymous. Bridge evidence: Static retrieval utility does not predict causal utility in multi-step agentic search. arXiv:2607.15253, 2026c.

Anonymous. Decision-aware memory cards. arXiv:2606.08151, 2026d.

Anonymous. Useful memories become faulty when continuously updated by LLMs. arXiv:2605.12978, 2026e.

Anonymous. HiMPO: Hindsight-informed memory policy optimization. arXiv:2606.16285, 2026f.

- Anonymous. Mem2ActBench: Evaluating memory-to-action consistency in LLM agents. arXiv:2601.19935, 2026g.
- Anonymous. MemAudit: Post-hoc auditing of poisoned agent memory via causal attribution. arXiv:2605.23723, 2026h.
- Anonymous. MemDelta: Controlled baselines and hidden confounds in agent memory evaluation. arXiv:2606.29914, 2026i.
- Anonymous. Memory transfer learning: How memories are transferred across domains in coding agents. arXiv:2604.14004, 2026j.
- Anonymous. MemTX: Transactional belief commit for agent memory. arXiv:2607.23929, 2026k.
- Anonymous. More skills, worse agents? when more skills are dangerous for LLM agents. arXiv:2605.24050, 2026l.
- Anonymous. Learning when to remember: Risk-sensitive abstention for agent memory retrieval. arXiv:2604.27283, 2026m.
- Anonymous. Grounding agent memory in contextual intent. arXiv:2601.10702, 2026n. *ACL 2026 Findings*.
- Chenao Cao et al. Remember me, refine me: A dynamic procedural memory framework for experience-driven agent evolution. In *ACL 2026 Findings*, 2026.
- Runsen Fang et al. Memp: Exploring agent procedural memory. arXiv:2508.06433, 2025.
- Zhaoxiang Feng, Mingyang Yao, and David Scott Lewis. When memories collide: Associative interference dynamics in lifelong agent memory. *ICLR 2026 Workshop NFAM*, 2026a.
- Zhaoxiang Feng, Mingyang Yao, and David Scott Lewis. Memory transplants for LLM agents: Disentangling architecture and content transfer under a code-to-math shift. *ICLR 2026 Workshop on Memory Agents (MemAgents)*, 2026b.
- Zhaoxiang Feng, Mingyang Yao, Charlie Sun, and David Scott Lewis. Which memory operation drives recovery? a factorial study of retrieve, write, and manage adaptation under domain shift. *ICLR 2026 Workshop on Lifelong Agents (LLA)*, 2026c.
- Anonymous Jalaldoust and Elias Bareinboim. Transportable representations for domain generalization. In *AAAI*, 2024.
- Ming Liang et al. Learning how to remember: A meta-cognitive management method for structured and transferable agent memory. In *ACL 2026 Findings*, 2026.
- Alessio Rocchi. Does memory credit travel? paired factorial audits of LLM-agent memory. SSRN 7160321, 2026. doi:10.2139/ssrn.7160321.
- Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. In *NeurIPS*, 2023.
- Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté, Yonatan Bisk, Adam Trischler, and Matthew Hausknecht. ALFWorld: Aligning text and embodied environments for interactive learning. *ICLR*, 2021.
- Qwen Team. Qwen2.5 technical report. arXiv:2412.15115, 2024.
- Ziqi Wang et al. Agent workflow memory. In *ICML*, 2025.
- Yifan Wu, Roshan Shariff, Tor Lattimore, and Csaba Szepesvári. Conservative bandits. In *ICML*, 2016.
- Guilin Zhang et al. Adaptive memory admission control. arXiv:2603.04549, 2026.
- Yan Zhang and Shibo Li. ConsistencyGate: Pruning LLM memories through self-consistency. arXiv:2607.22962, 2026.
- Andrew Zhao, Daniel Huang, Quentin Xu, Matthieu Lin, Yong-Jin Liu, and Gao Huang. ExpeL: LLM agents are experiential learners. In *AAAI*, 2024.

A H-DC PROTOCOL, PRE-ANALYSIS AMENDMENTS, AND AUDIT

Design and analysis of the deployment test of §7 follow a timestamped pre-analysis plan (GATE.PROTOCOL Part IV, amended by Part IV-A), under the project’s governance rule: implementation gaps were first self-analyzed, then adjudicated with an external reviewer, then registered. We disclose the provenance of these timestamps rather than round it to “pre-registered”: repository history shows the card builder was implemented (16:59 JST) and outcome-bearing rollouts had begun ($\approx$ 17:01 JST) before the plan text was committed (17:06 JST), and the plan text precedes the analysis-code and results commits (17:10 and 18:10 JST). That no outcomes were inspected before the amendments is author-reported, not independently verifiable from the artifacts; the plan is exactly what the timestamps establish — an analysis plan frozen during construction and before analysis. Amendments, all registered before the primary analyses were run:

1. **3B control arm added.** The registration had assumed all controls already existed; the naive-truncation arm in fact existed only for 7B. The 3B RAW arm (2,560 rollouts, identical cards/harness/seeds) was run as a registered correction, keeping E1 estimable at both scales and the Holm family ($m = 6$) intact.
2. **Paired-budget control made primary.** Realized DSLICE card length (mean 237.5 tokens) is genuinely shorter than the 300-token cap, so a “shorter is conservative” argument would not have sufficed: prompt length itself can affect success. The RAW_MATCHED control—same transcript truncated to the *exact* paired token count (per-card $|\Delta| = 0, 640/640$)—became the primary contrast; RAW(300) is reported as the secondary senior baseline.
3. **Packaging disclosed.** Action copies are canonical dumps of the parsed action JSON (strips harness parse artifacts; 18/4,510 affected steps, with verbatim fallback wherever a tail could carry recoverable content); exact-duplicate non-mutating steps are deduplicated; the escalation ladder reaches stages $\{0, 1, 2, 3, 4, 5\} = \{14, 76, 281, 8, 252, 9\}$ cards. Claims therefore attach to the DSLICE *package*, not to any single rule.

Execution-environment checks, archived in `pilot/dslice/A6.CHECKS.json`: a full metadata scan over every rollout row finds exactly one config hash and one vLLM/torch/GPU environment version per arm (the 3B pilot N/Q rows differ only in a GPU-memory-utilization resource field, unrelated to tasks, decoding, or prompts), with zero rows missing any audited field; the one post-hoc harness change (single-BOS tokenization, introduced for Llama) is verified token-id-identical on both Qwen2.5 sizes (pinned 3B/7B tokenizer snapshots) across first-turn and multi-turn prompt shapes with and without memory, i.e., a no-op for every arm it touches. Inference: aligned family-cluster bootstrap (2000 reps) with the primary Holm family $m = 6$ $\{E1 \times 2$ models, $E2 \times 2$, HFR-noninferiority $\times 2\}$; trap noninferiority uses one-sided bootstrap bounds against the $\pm 5pp$ margins. Full numbers: `pilot/dslice/HDC.RESULTS.json`; protocol text: `GATE.PROTOCOL.md` Part IV/IV-A.

Provenance-stratified sensitivity (post-hoc). 188 of the 640 memories (29.4%; A00/A01/A10/A11 = 52/49/43/44) are deterministic reconstructions of the evaluator’s oracle plan (harvest fallback), not model-generated episodes (§7.1). The E1 estimand uses the A10/A11 cells, i.e. 233 of the 452 model-generated cards and 87 of the 188 oracle-reconstructed cards. Re-running the aligned family-cluster bootstrap (2000 reps, seed 1234; outside the Holm family) within strata gives $E1 = +0.041 [-0.078, +0.149]$ at 7B and $-0.226 [-0.347, -0.104]$ at 3B on the model-generated stratum (36 families), versus $+0.424 [+0.255, +0.615]$ at 7B and $+0.297 [+0.167, +0.473]$ at 3B on the oracle-reconstructed stratum (20 families); the pooled post-hoc cross-model E1 difference is $+0.230 [+0.131, +0.338]$. The direct provenance interaction (one aligned family-cluster bootstrap over the union family universe—40 families, 16 of which occur in both strata—so the stratum E1s themselves are preserved and shared-family covariance is retained) is $+0.382 [+0.166, +0.606]$ at 7B and $+0.523 [+0.306, +0.783]$ at 3B (oracle minus authentic). Full numbers: `pilot/dslice/HDC.SENSITIVITY.FALLBACK.json`.

B REPRODUCIBILITY

All artifacts at github.com/anonymous/CausalMemBench upon de-anonymization:

- generator `pilot/generate_families.py` (pilot seed 20260807, H3 seed 20260809); sealed-vs-public split; oracle walker execution logs confirming 800/800 and 640/640 legal terminals;
- environment `pilot/env_relationalops.py`, `DSLpilot/program_dsl.py`, `harnesspilot/harness.py` (vLLM 0.6.6.post1 / torch 2.5.1+cu124 on RTX A5000 24GB, fp16, temperature 0.7, top-p 0.9, max 12 steps $\times$ 512 tokens);
- rollouts JSONL: 41,216 rows (7,680 pilot + 7,680 H-C + 9,216 H3 + 3,840 Llama-8B replication + 12,800 H-DC deployment grids) with per-row commit, config-hash, model, versions, GPU;
- protocols: GATE_PROTOCOL Parts I–III pre-registered (2026-08-07, before implementation and analysis); Part IV/IV-A is a timestamped pre-analysis plan frozen during H-DC construction and before its analysis-code/result commits (provenance disclosed in Appendix A);
- audit suite `pilot/audit/{probes, difficulty_robust, continuous_s, equivalence_audit}.py`, blind annotations, judge outputs verbatim; DSLICE/RAW_MATCHED card builders with per-card QA maps;
- analysis: family-cluster bootstrap utilities, Holm/TOST routines, and the exact `analyze.py/json` dumps reported here;
- design registers: RESEARCH_LEDGER (pre-registration chain + external-review outcomes), §13 artifacts, construction rulings for H3 cards.

C FAMILIES AND ARCHETYPES

40 pilot families = 4 archetypes $\times$ 2 domains (crm/inv/ticket/cal) $\times$ 5 instances; 32 H3 families = 4/archetype $\times$ 8 schemas with disjoint parameters. For each family: 4 program-preserving siblings + 1 near-miss + a10 cross-domain partner + a00 unrelated partner, style counterbalanced by a 6-template Latin square, memories 200–300 Qwen-1.5B tokens measured by tokenizer, similarity buckets frozen against pilot distribution anchors.

D SECOND-FAMILY REPLICATION (COMPLETED)

Llama-3.1-8B (NousResearch/Meta-Llama-3.1-8B-Instruct) replicates the identical grid (3840 rollouts, same config/harness). Success rates $N=.373$, $Q=.408$, $A00=.364$, $A01=.452$, $A10=.411$, $A11=.544$. Estimates: $\tau_{\text{struct}} = +0.047$ [$+0.020, +0.073$] (7B-like, opposite 3B); $\tau_{\text{trap}} = +0.087$ [$+0.043, +0.134$]; $\tau_{\text{replay}} = +0.133$ [$+0.069, +0.193$]; $\text{HFR} = 0.100$ [$0.075, 0.125$]; $\tau_{\text{context}} = +0.034$ [$+0.010, +0.058$] (family-specific: Qwen has ≈ 0). Replay-heavy decomposition: the A11 leg carries 79% of the matched effect, mirroring the 7B 72% share. See `pilot/llama8b/LLAMA8B.RESULTS.json`.

E EXTERNAL VALIDITY CHECK (ALFWORLD, COMPLETED)

A coarse effect-ordering check in ALFWorld (Shridhar et al., 2021) (Qwen2.5-7B; 3 task families — clean/, heat/, pick-two — $\times$ 4 siblings $\times$ 4 cells $\times$ 3 seeds, 144 rollouts, max 30 steps, env-verified win) maps pilot cells as $N/R(\approx A11)/S(\approx A10)/X(\approx A01$ with the critical step merely omitted, not actively flipped). Verdict, reported without overclaim: **partially reproduced**. Aggregate ordering $R(0.222) > S(0.194) > N(0.167)$ supports replay-superior-to-structural-superior-to-nothing, at $n=36/\text{cell}$ with overlapping Wilson intervals—direction only, and the clean family alone reproduces the full pilot pattern ($R=S=.33 > X=.08 > N=.00$). **The near-miss harm did not reproduce** ($X=0.250 \geq N$, paired flips 3 harmful vs 6 helpful): our ALFWorld near-miss is a recoverable omission rather than an active polarity/order contradiction, so a 7B model can recover the missing step from the visible goal text. Reading: replay-over-structure ordering survives a non-synthetic env, while harmful near-miss transfer appears contingent on whether the trap actively contradicts the task rather than omits a step. Full construction and deviation log: `pilot/external/EXTERNAL_VALIDATION.md`.